

Capable LLMs converge on the same wrong answer at rates that stay nearly flat as ensemble size grows, while independence chance drops geometrically — a 20× separation at k=5.

