

*A simple routing rule — flag questions where mid-tier LLMs unanimously agree but the frontier model disagrees — catches a meaningful fraction of silent ensemble failures at low human-review cost.*

